

# ENM53 I: Data-driven modeling and probabilistic scientific computing

## *Lecture #23: Generative adversarial networks*

Paris Perdikaris  
April 28, 2020



# Tricks of the trade

- Variational bounds
- Density re-parametrizations
- Density ratio estimation
- Variational optimization/evolution strategies
- Adversarial games

# The density ratio trick

The central task in the above five statistical quantities is to efficiently compute the ratio  $r(x)$ . In simple problems, we can compute the numerator and the denominator separately, and then compute their ratio. Direct estimation like this will not often be possible: each part of the ratio may itself involve intractable integrals; we will often deal with high-dimensional quantities; and we may only have samples drawn from the two distributions, not their analytical forms.

This is where the *density ratio trick* or formally, *density ratio estimation*, enters: it tells us to construct a binary classifier  $S(x)$  that distinguishes between samples from the two distributions. We can then **compute the density ratio using the probability given by this classifier**:

$$r(x) = \frac{\rho(x)}{q(x)} = \frac{S(x)}{1 - S(x)}$$

To show this, imagine creating a data set of  $2N$  elements consisting of pairs (data  $x$ , label  $y$ ):

- $N$  data points are drawn from the distribution  $\rho$  and assigned a label  $+1$ .
- The remaining  $N$  data points are drawn from distribution  $q$  and assigned label  $-1$ .


# Density ratio estimation by probabilistic classification

$$\text{KL}[p(\mathbf{x})||q(\mathbf{x})] := \int \log \frac{p(\mathbf{x})}{q(\mathbf{x})} p(\mathbf{x}) d\mathbf{x} = \mathbb{E}_{p(\mathbf{x})} \left[ \log \frac{p(\mathbf{x})}{q(\mathbf{x})} \right]$$

Estimating density ratios is a challenging task:


- Each part of the ratio may itself involve intractable integrals
- We often deal with high-dimensional quantities.
- We may only have samples drawn from the two distributions, not their analytical forms.

This is where the **density ratio trick** enters:  
it allows us to construct a binary classifier  
that distinguishes between samples from the  
two distributions.



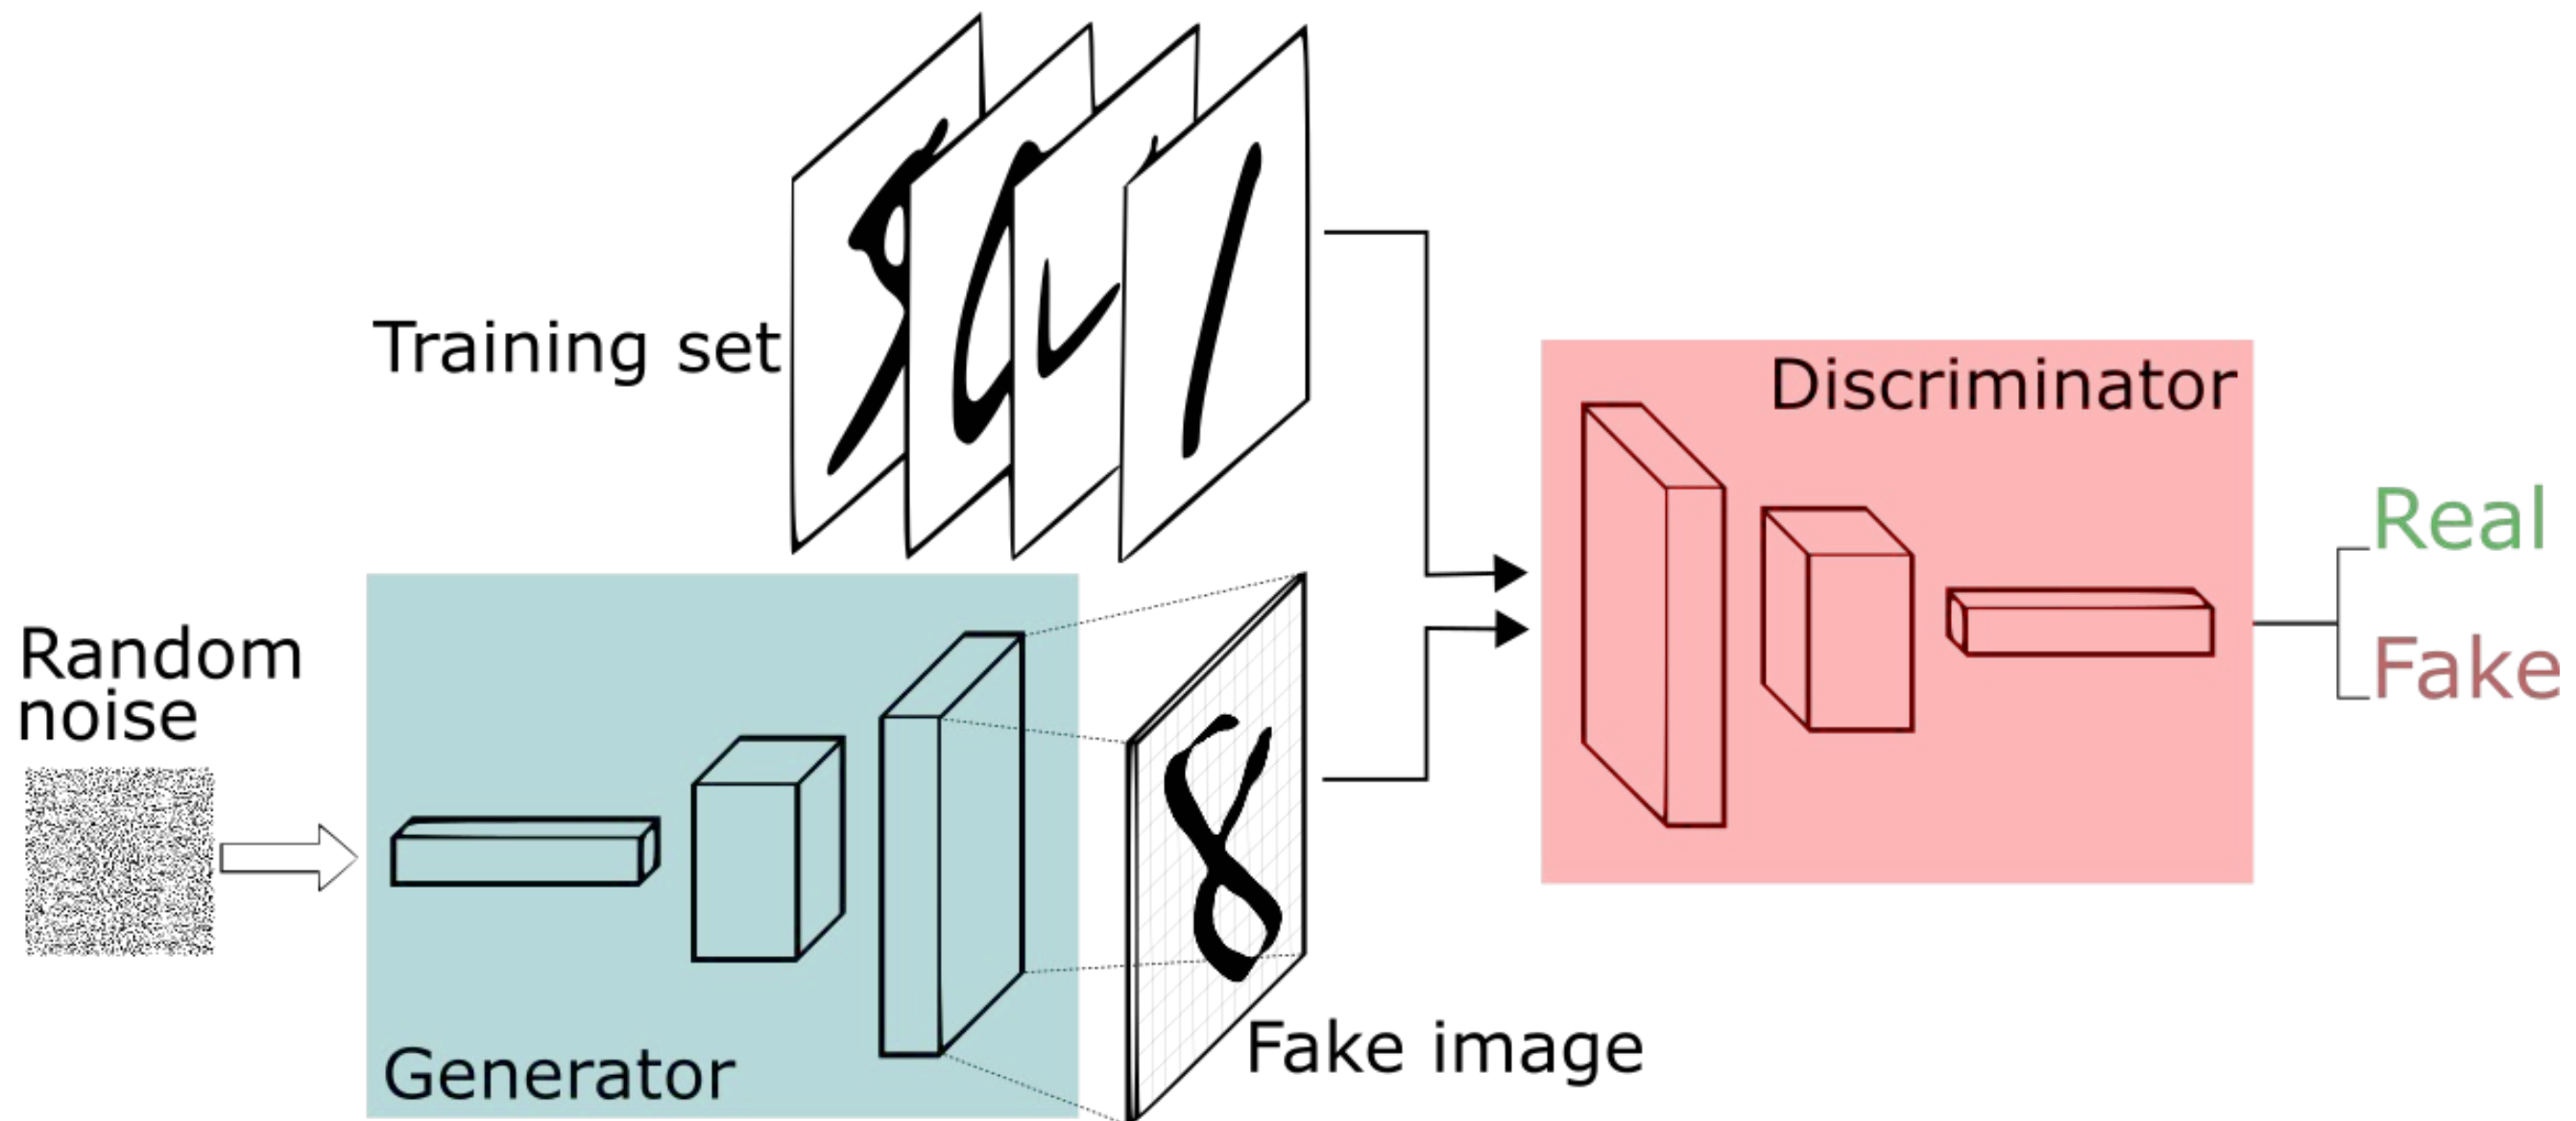
*The density ratio gives the correction factor  
needed to make two distributions equal.*

$$\begin{aligned} r(x) &= \frac{\rho(x)}{q(x)} = \frac{p(x|y = +1)}{p(x|y = -1)} \\ &= \frac{p(y = +1|x)p(x)}{p(y = +1)} \bigg/ \frac{p(y = -1|x)p(x)}{p(y = -1)} \\ &= \frac{p(y = +1|x)}{p(y = -1|x)} = \frac{p(y = +1|x)}{1 - p(y = +1|x)} = \frac{\mathcal{S}(x)}{1 - \mathcal{S}(x)} \end{aligned}$$





# Generative adversarial networks



$$\min_G \max_D V(D, G) = \mathbb{E}_{q(\mathbf{x})} [\log(D(\mathbf{x}))] + \mathbb{E}_{p(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$


# Generative adversarial networks

## PROGRESSIVE GROWING OF GANs FOR IMPROVED QUALITY, STABILITY, AND VARIATION

Submitted to ICLR 2018

# Conditional generative adversarial networks

Mirza and Osindero (2014)



**GAN**  $\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$

**CGAN**  $\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\underline{\mathbf{y}})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\underline{\mathbf{y}})))]$

# Conditional generative adversarial networks


## High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs

Ting-Chun Wang<sup>1</sup>, Ming-Yu Liu<sup>1</sup>, Jun-Yan Zhu<sup>2</sup>, Andrew Tao<sup>1</sup>,  
Jan Kautz<sup>1</sup>, Bryan Catanzaro<sup>1</sup>

<sup>1</sup>NVIDIA Corporation   <sup>2</sup>University of California, Berkeley



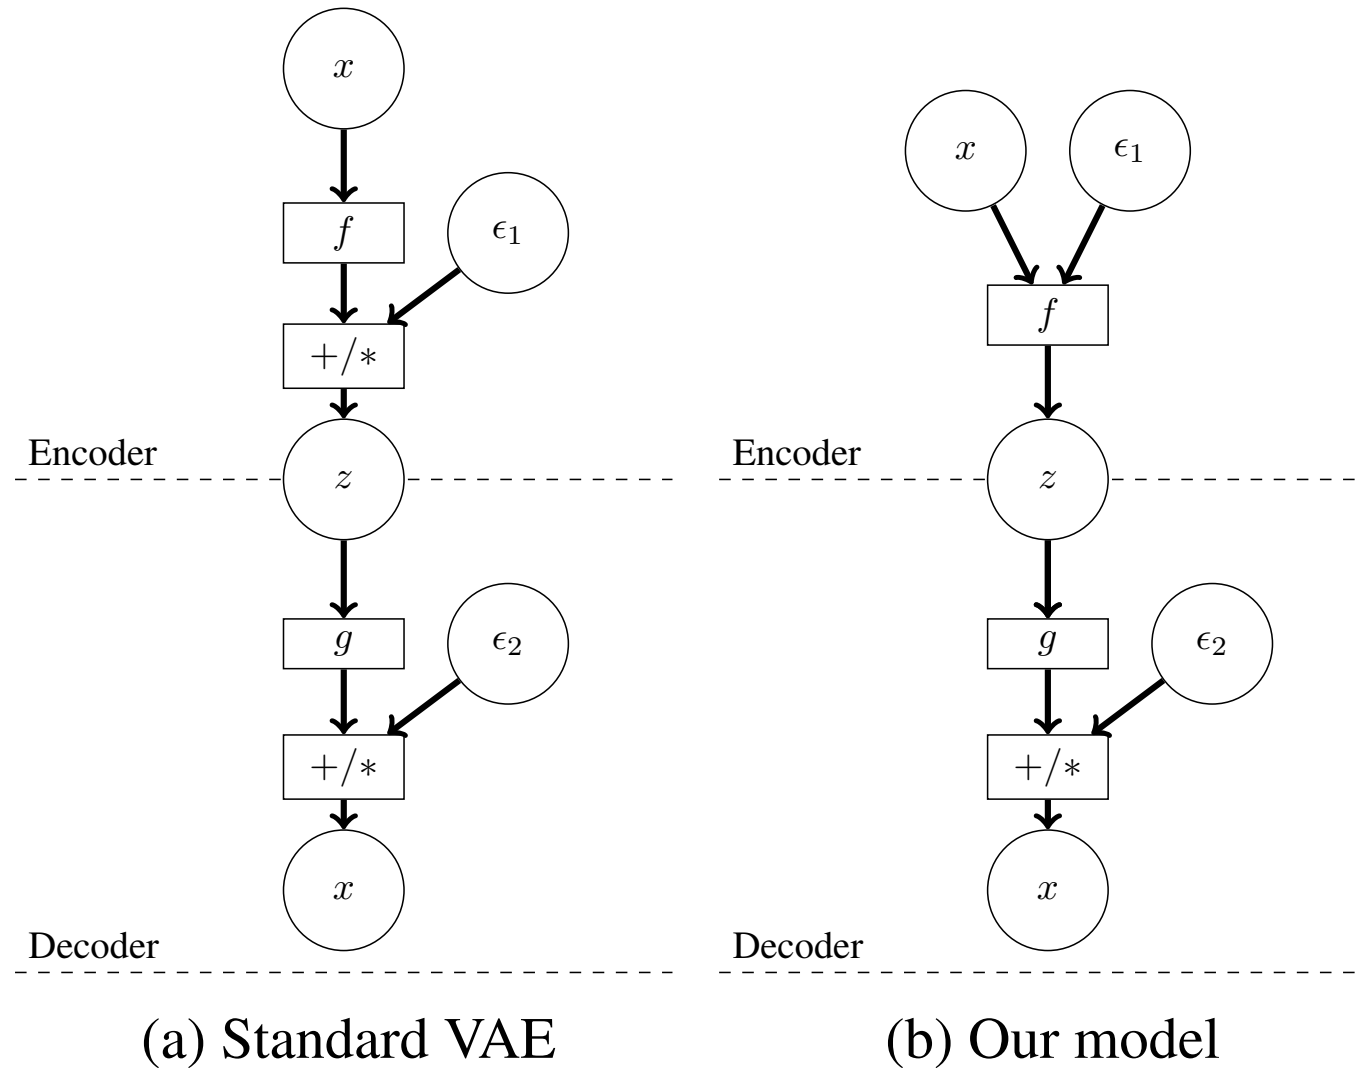

# Adversarially learned inference



$$\min_G \max_D V(D, G) = \mathbb{E}_{q(\mathbf{x})} [\log(D(\mathbf{x}, G_z(\mathbf{x})))] + \mathbb{E}_{p(\mathbf{z})} [\log(1 - D(G_x(\mathbf{z}), \mathbf{z}))]$$

# Adversarial Variational Bayes

Instead of using approximating distributions of a given pre-defined form (e.g. Gaussian) we can implicitly parametrize them using deep neural networks.



## Algorithm 1 Adversarial Variational Bayes (AVB)

- 1:  $i \leftarrow 0$
- 2: **while** not converged **do**
- 3:   Sample  $\{x^{(1)}, \dots, x^{(m)}\}$  from data distrib.  $p_{\mathcal{D}}(x)$
- 4:   Sample  $\{z^{(1)}, \dots, z^{(m)}\}$  from prior  $p(z)$
- 5:   Sample  $\{\epsilon^{(1)}, \dots, \epsilon^{(m)}\}$  from  $\mathcal{N}(0, 1)$
- 6:   Compute  $\theta$ -gradient (eq. 3.7):  

$$g_{\theta} \leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\theta} \log p_{\theta} (x^{(k)} \mid z_{\phi} (x^{(k)}, \epsilon^{(k)}))$$
- 7:   Compute  $\phi$ -gradient (eq. 3.7):  

$$g_{\phi} \leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\phi} [-T_{\psi} (x^{(k)}, z_{\phi}(x^{(k)}, \epsilon^{(k)})) + \log p_{\theta} (x^{(k)} \mid z_{\phi}(x^{(k)}, \epsilon^{(k)}))]$$
- 8:   Compute  $\psi$ -gradient (eq. 3.3) :  

$$g_{\psi} \leftarrow \frac{1}{m} \sum_{k=1}^m \nabla_{\psi} [\log (\sigma(T_{\psi}(x^{(k)}, z_{\phi}(x^{(k)}, \epsilon^{(k)})))) + \log (1 - \sigma(T_{\psi}(x^{(k)}, z^{(k)})))]$$
- 9:   Perform SGD-updates for  $\theta$ ,  $\phi$  and  $\psi$ :  

$$\theta \leftarrow \theta + h_i g_{\theta}, \quad \phi \leftarrow \phi + h_i g_{\phi}, \quad \psi \leftarrow \psi + h_i g_{\psi}$$
- 10:    $i \leftarrow i + 1$
- 11: **end while**

Huszár, F. (2017). Variational inference using implicit distributions. arXiv preprint arXiv:1702.08235.

Mescheder, L., Nowozin, S., & Geiger, A. (2017). Adversarial variational Bayes: Unifying variational autoencoders and generative adversarial networks. arXiv preprint arXiv:1701.04722.

Makhzani, A., Shlens, J., Jaitly, N., Goodfellow, I., & Frey, B. (2015). Adversarial autoencoders. arXiv preprint arXiv:1511.05644.